



★ Member-only story


When and How to Create Indexes in a SQL Database

Nagarjun (Arjun) Nagesh · [Follow](#)

3 min read · Aug 14, 2024



1



Indexes are critical components of relational databases, designed to enhance the speed and efficiency of data retrieval operations. By organizing data in a manner that facilitates quicker searches, indexes can significantly reduce query execution time, especially in large datasets. However, improper use or overuse of indexes can degrade performance. This article delves into the ideal scenarios for creating indexes, how to create them using SQL, and their impact on database performance.

What is an Index in SQL?

An index in SQL is a database object that improves the speed of data retrieval operations on a table. Think of an index like an index in a book: it allows the database engine to locate rows without scanning the entire table.

How Indexes Work

When a query is executed, the database can use the index to quickly find the relevant rows, similar to how you use an index in a book to locate a specific topic. Without an index, the database might need to scan every row to find the data, which is much slower.

Ideal Scenarios for Creating Indexes

1. Frequent Searches on Specific Columns

- **Scenario:** Columns that are frequently used in `WHERE` clauses, `JOIN` conditions, or as part of a `SELECT` query.
- **Example:** A large customer database frequently queried by `customer_id`.

SQL Statement:

```
CREATE INDEX idx_customer_id ON customers(customer_id);
```

2. Large Tables with Heavy Read Operations

- **Scenario:** Tables with a vast number of records where read operations (`SELECT` queries) are more common than write operations.
- **Example:** A transaction history table where you often search by `transaction_date`.

SQL Statement:

```
CREATE INDEX idx_transaction_date ON transactions(transaction_date);
```

3. Columns Used in JOINS

- **Scenario:** Columns frequently used to join tables in SQL queries. Indexing these columns can make join operations faster.
- **Example:** An `orders` table and a `customers` table frequently joined on `customer_id`.

SQL Statement:

```
CREATE INDEX idx_orders_customer_id ON orders(customer_id);
```

4. Unique Constraints

- **Scenario:** Columns that must maintain unique values. Indexes can enforce uniqueness and improve lookup efficiency.
- **Example:** An email column in a user table where each email must be unique.

SQL Statement:

```
CREATE UNIQUE INDEX idx_unique_email ON users(email);
```

5. Composite Indexes

- **Scenario:** When queries often filter or sort by multiple columns. A composite index can be more efficient than multiple single-column indexes.
- **Example:** A query that frequently searches for orders by `customer_id` and `order_date`.

SQL Statement:

```
CREATE INDEX idx_customer_date ON orders(customer_id, order_date);
```

When Not to Use Indexes

- **Tables with Heavy Write Operations:** Indexes slow down `INSERT`, `UPDATE`, and `DELETE` operations because the index needs to be updated whenever data is modified.
- **Small Tables:** For small tables, the overhead of maintaining an index may outweigh its benefits.
- **Columns with Low Selectivity:** Indexing a column with many duplicate values (like a boolean field) is often not beneficial.

Impact of Indexes on Performance

- **Improved Query Performance:** Indexes significantly speed up data retrieval by reducing the number of rows the database engine needs to scan.
- **Slower Write Operations:** Every time data is inserted, updated, or deleted, the index must also be updated, which can slow down these operations.
- **Increased Storage Requirements:** Indexes consume additional disk space.

Monitoring and Maintaining Indexes

- **Use EXPLAIN PLAN:** To analyze how a query is executed and whether an index is being used effectively.
- **Regularly Rebuild or Reorganize Indexes:** To maintain performance, especially in databases with frequent data modifications.

Practical Example: Optimizing a Query with Indexes

Consider a scenario where you have a large `orders` table, and you're frequently running the following query:

```
SELECT * FROM orders WHERE customer_id = 12345 AND order_date = '2024-01-01';
```

Without an index, this query might require a full table scan. By creating a composite index, you can dramatically improve performance:

```
CREATE INDEX idx_customer_date ON orders(customer_id, order_date);
```

Conclusion

Indexes are powerful tools in SQL databases that can drastically improve query performance when used appropriately. Understanding when and how to create indexes can lead to significant performance gains, especially in large databases or systems with complex queries. However, it's essential to strike a balance, as unnecessary or excessive indexing can degrade performance.

By regularly monitoring query performance and adjusting indexes accordingly, you can ensure that your database remains efficient and responsive.

References

- **Official MySQL Documentation on Indexes:** [MySQL Indexes](#)
- **PostgreSQL Documentation on Indexes:** [PostgreSQL Indexes](#)
- **SQL Server Index Design Guide:** [SQL Server Indexing](#)

These references provide additional insights and best practices for implementing and maintaining indexes in various SQL-based databases.